



Variance partitioning: hidden vs surprisal


gpt2




opt-125m




pythia-160m



pythia-410m



qwen2.5-0.5b



qwen2.5-0.5b-instruct

